

Final Project

Predicting Wine Quality Based on Chemical Properties

1. Introduction

Objective:

- The goal of this project is to predict wine quality using chemical attributes such as acidity, alcohol content, and sugar levels. Additionally, actionable insights are derived to optimize production and marketing strategies.

Dataset:

- Name: Wine Quality Dataset.
- Size: 1,999 rows and 12 columns.
- Features: Chemical attributes such as fixed acidity, volatile acidity, residual sugar, and the target variable quality (rated 0–10).

Workflow:

1. Exploratory Data Analysis (EDA)
2. Data Preprocessing
3. Feature Engineering
4. Model Building and Evaluation
5. Insights and Recommendations

2. Exploratory Data Analysis (EDA)

Purpose: To identify patterns, correlations, and potential issues in the dataset.

Key Observations

1. Correlation:

- Alcohol content shows a strong positive correlation with wine quality.
- Sulphates and acidity also influence quality significantly.

2. Distribution of Quality:

- Most wines are rated as medium quality (5 or 6), indicating an imbalanced target variable.

3. Outliers:

- Residual sugar and volatile acidity contain outliers.

Code Snippet:

```
*[7]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

file_path = 'winequality-dataset_updated.csv' # Replace with your dataset path
wine_data = pd.read_csv(file_path)

print("Dataset Info:")
print(wine_data.info())

print("\nBasic Statistics:")
print(wine_data.describe())

print("\nMissing Values:")
print(wine_data.isnull().sum())

plt.figure(figsize=(10, 8))
sns.heatmap(wine_data.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Heatmap")
plt.show()

plt.figure(figsize=(8, 6))
sns.countplot(x='quality', data=wine_data, palette='viridis')
plt.title("Distribution of Wine Quality")
plt.xlabel("Wine Quality")
plt.ylabel("Count")
plt.show()

plt.figure(figsize=(10, 6))
sns.boxplot(data=wine_data[['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar']])
plt.title("Boxplots of Key Features")
plt.xticks(rotation=45)
plt.show()

selected_features = ['alcohol', 'pH', 'sulphates', 'quality']
sns.pairplot(wine_data[selected_features], hue='quality', palette='husl')
plt.show()
```

```

plt.figure(figsize=(8, 6))
sns.barplot(x='quality', y='alcohol', data=wine_data, palette='magma')
plt.title("Alcohol Content vs Wine Quality")
plt.xlabel("Wine Quality")
plt.ylabel("Alcohol Content")
plt.show()

plt.figure(figsize=(8, 6))
sns.boxplot(x='quality', y='pH', data=wine_data, palette='coolwarm')
plt.title("pH vs Wine Quality")
plt.xlabel("Wine Quality")
plt.ylabel("pH")
plt.show()

```

Dataset Info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1999 entries, 0 to 1998
Data columns (total 12 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   fixed acidity      1999 non-null   float64
 1   volatile acidity   1999 non-null   float64
 2   citric acid        1999 non-null   float64
 3   residual sugar     1999 non-null   float64
 4   chlorides          1999 non-null   float64
 5   free sulfur dioxide 1999 non-null   float64
 6   total sulfur dioxide 1999 non-null   float64
 7   density            1999 non-null   float64
 8   pH                 1999 non-null   float64
 9   sulphates          1999 non-null   float64
 10  alcohol            1999 non-null   float64
 11  quality             1999 non-null   int64 

```

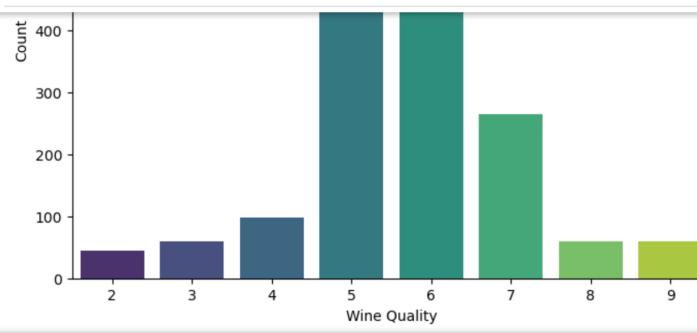
```

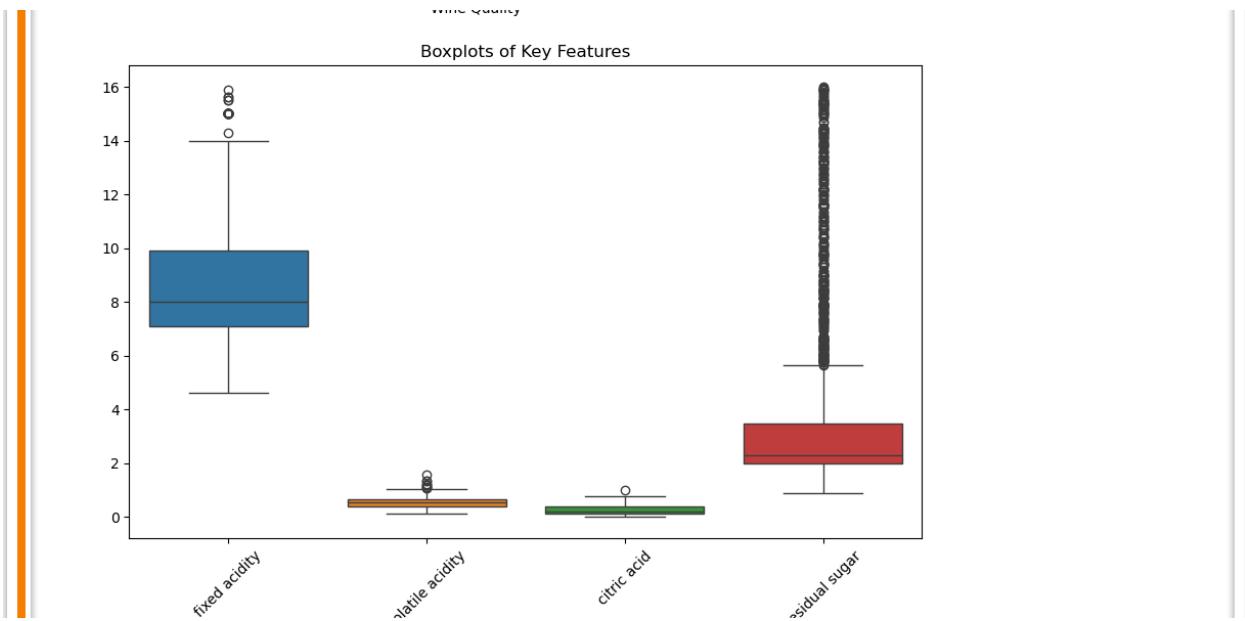
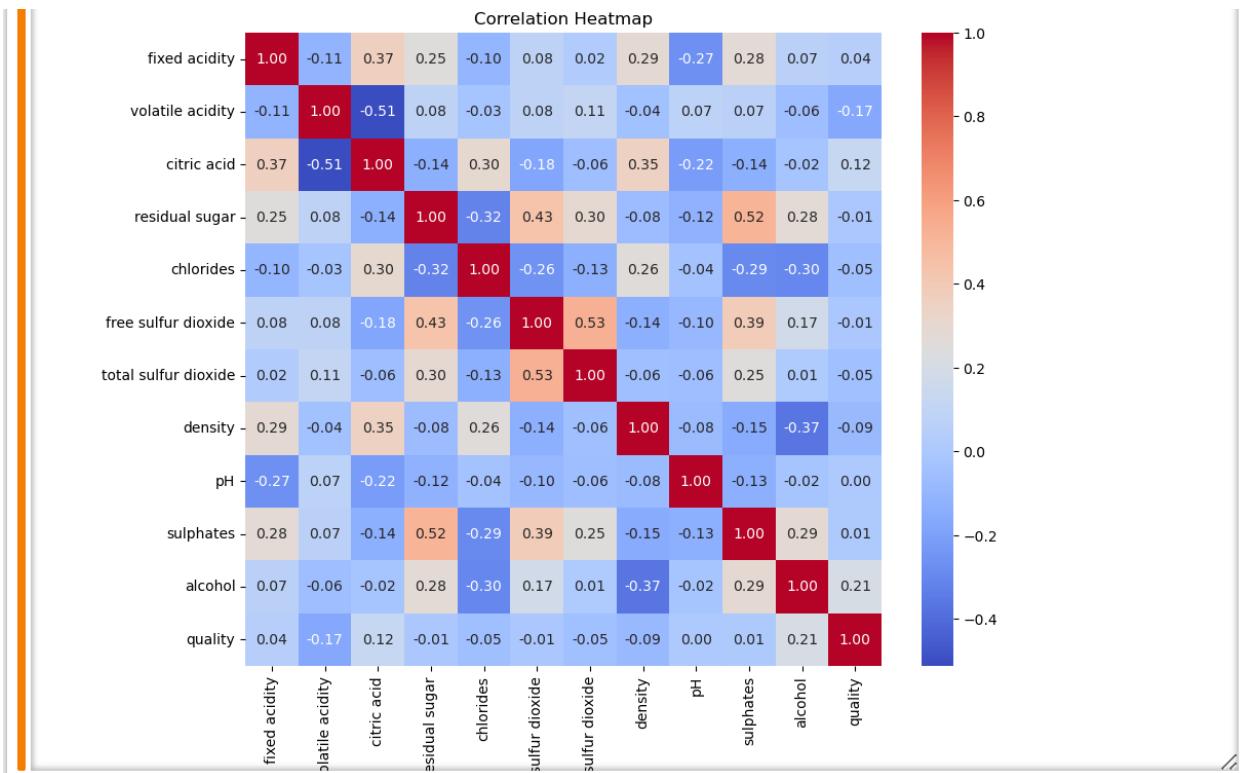
dtypes: float64(11), int64(1)
memory usage: 187.5 KB
None

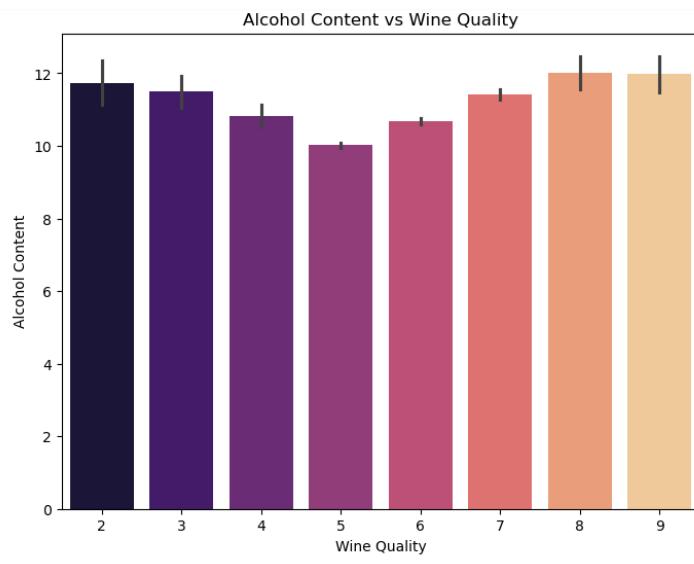
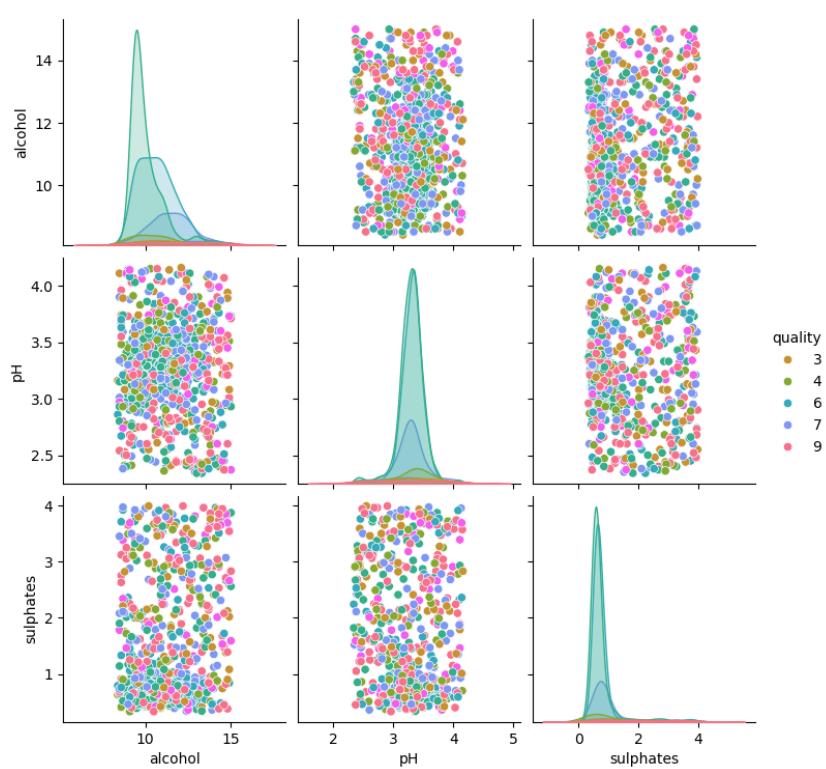
Basic Statistics:
fixed acidity  volatile acidity  citric acid  residual sugar \
count    1999.000000    1999.000000  1999.000000  1999.000000
mean      8.670335    0.541773   0.246668   3.699090
std       2.240023    0.180381   0.181348   3.290201
min       4.600000    0.120000   0.000000   0.900000
25%      7.100000    0.400000   0.110000   2.000000
50%      8.000000    0.530000   0.200000   2.300000
75%      9.900000    0.660000   0.385000   3.460000
max     15.900000    1.580000   1.000000  15.990000

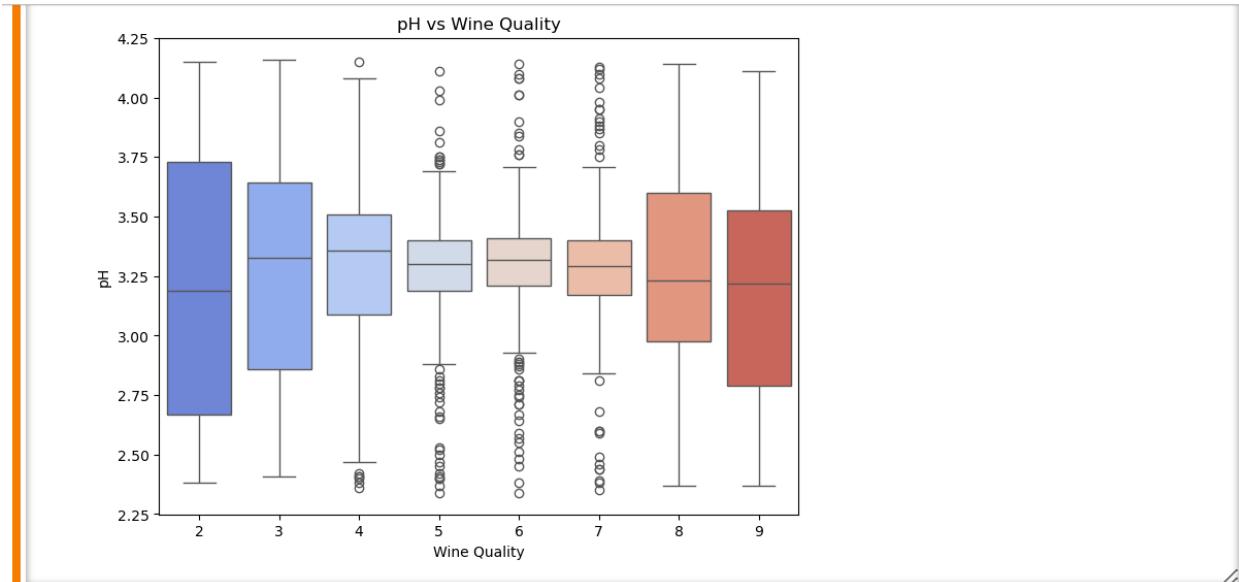
chlorides  free sulfur dioxide  total sulfur dioxide  density \
count    1999.000000    1999.000000  1999.000000  1999.000000
mean      0.075858    20.191096    52.617809   0.996477

```









Visualization

- **Correlation Heatmap:** Highlights relationships between features.
- **Quality Distribution Plot:** Shows the imbalance in the target variable.

3. Data Preprocessing

Purpose: To clean and standardize the dataset for modeling.

Steps Taken

1. **Missing Values:**
 - No missing values were detected in the dataset.
2. **Outlier Treatment:**
 - Outliers in residual sugar and volatile acidity were capped using the IQR method.
3. **Feature Scaling:**
 - Standard Scaler was applied to normalize continuous features.

Code Snippet:

```
def treat_outliers(column):
    Q1 = wine_data[column].quantile(0.25)
    Q3 = wine_data[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    wine_data[column] = np.where(wine_data[column] < lower_bound, lower_bound,
                                 np.where(wine_data[column] > upper_bound, upper_bound, wine_data[column]))

columns_to_check = ['fixed acidity', 'volatile acidity', 'citric acid',
                     'residual sugar', 'chlorides', 'free sulfur dioxide',
                     'total sulfur dioxide', 'density', 'pH', 'sulphates', 'alcohol']
for column in columns_to_check:
    treat_outliers(column)

scaler = StandardScaler()
scaled_features = scaler.fit_transform(wine_data[columns_to_check])

scaled_data = pd.DataFrame(scaled_features, columns=columns_to_check)

scaled_data['quality'] = wine_data['quality']

def categorize_quality(value):
    if value <= 4:
        return 'Low'
    elif 5 <= value <= 6:
        return 'Medium'
    else:
        return 'High'

X = scaled_data.drop('quality', axis=1)
y = scaled_data['quality']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

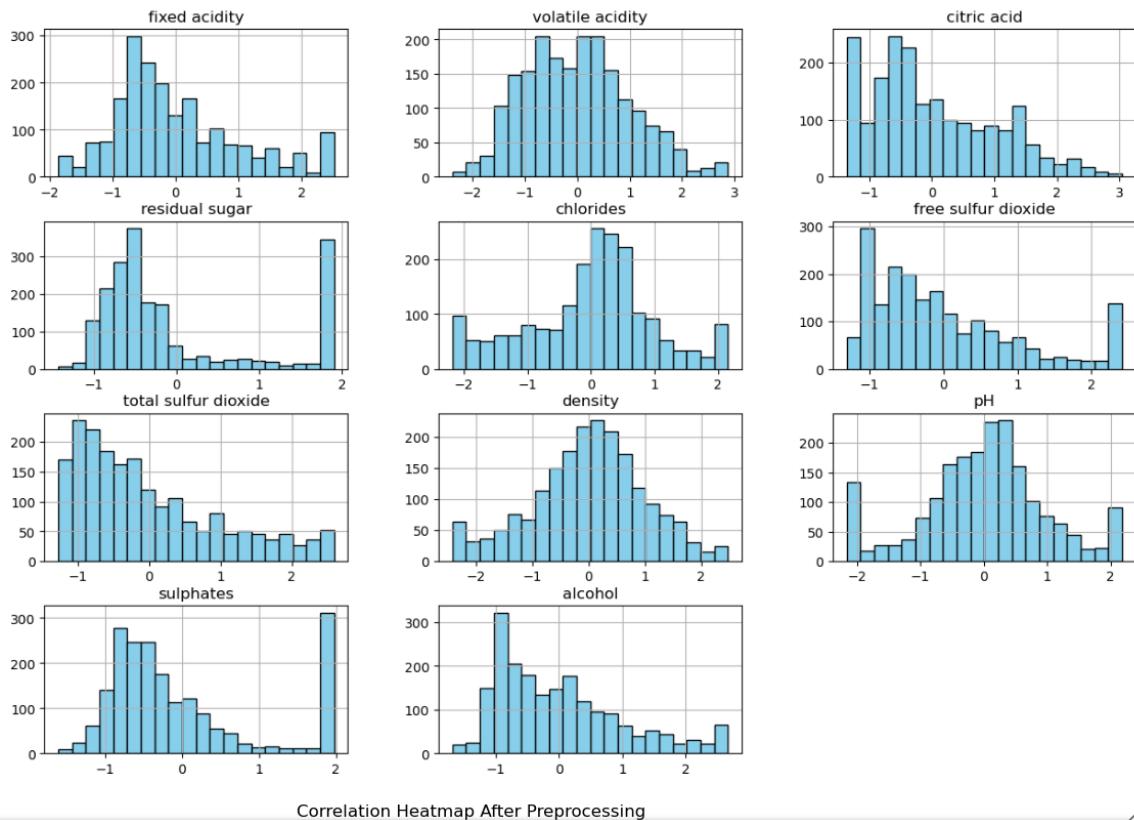
print("\nTraining Set Shape:", X_train.shape, y_train.shape)
print("Testing Set Shape:", X_test.shape, y_test.shape)

Missing Values Before Handling:
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                 0
sulphates          0
alcohol            0
quality            0
dtype: int64

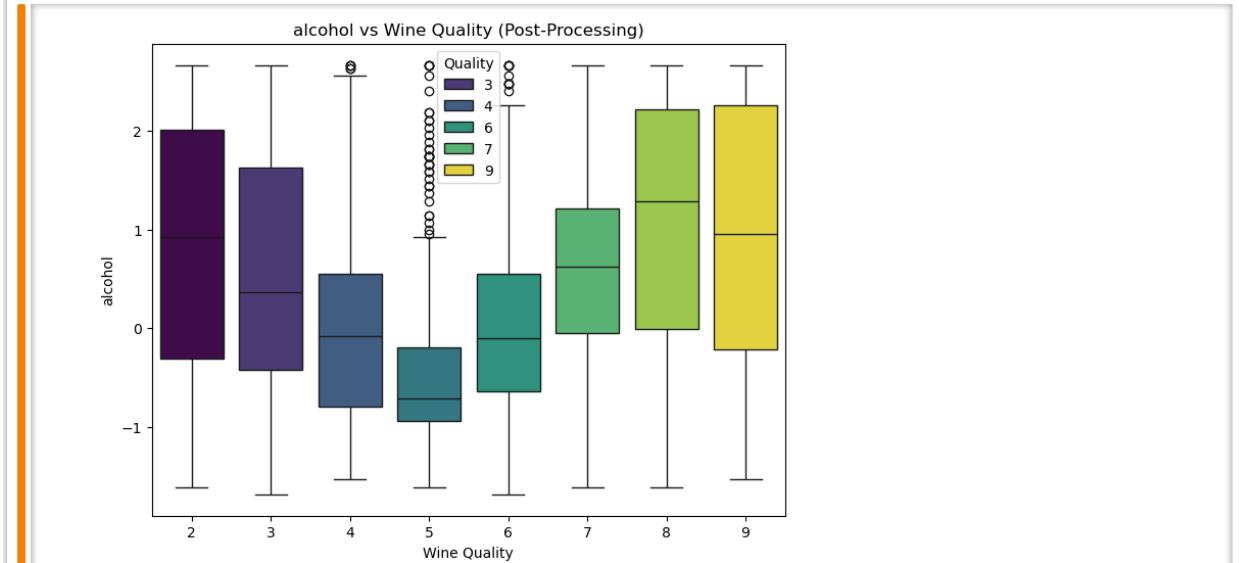
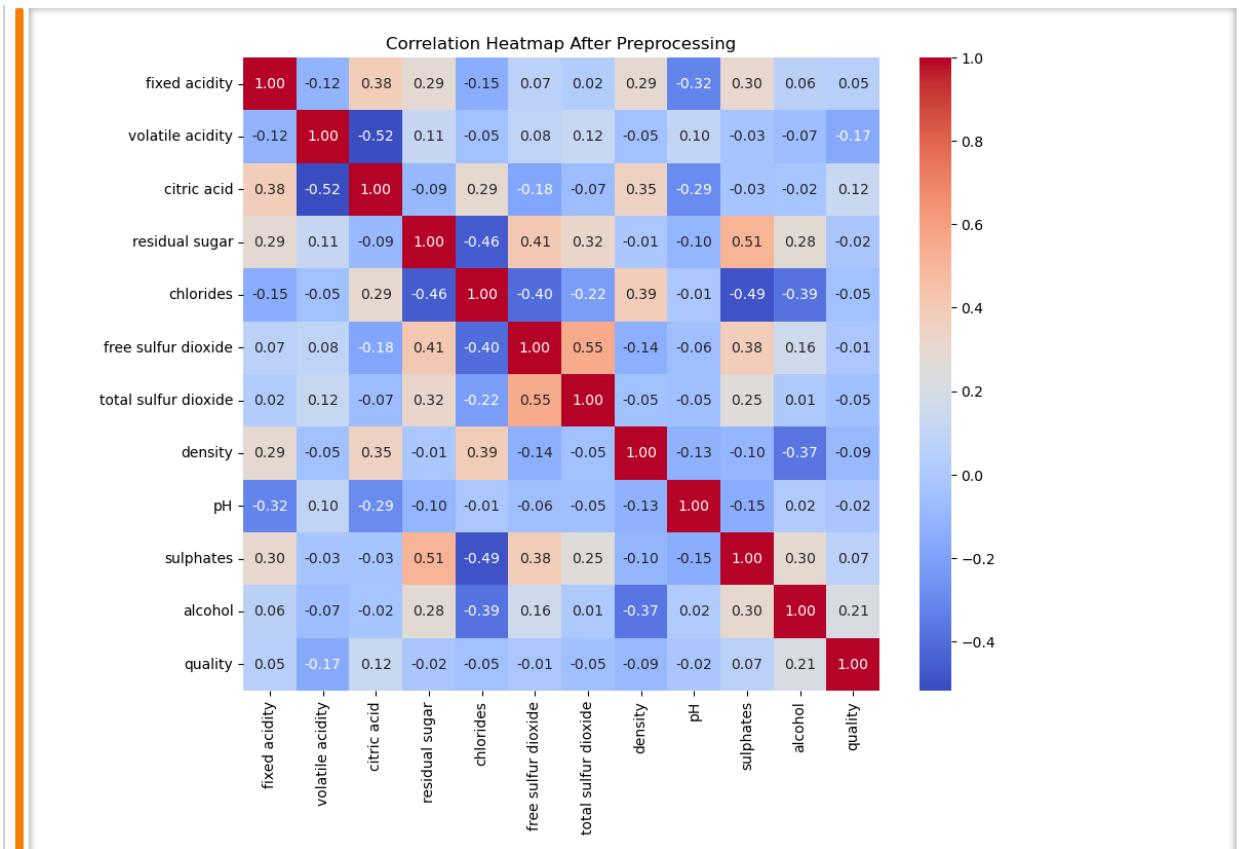
Missing Values After Handling:
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                 0
sulphates          0
alcohol            0
quality            0
dtype: int64

Training Set Shape: (1599, 11) (1599,)
Testing Set Shape: (400, 11) (400,)
```

Distribution of Features After Scaling



Correlation Heatmap After Preprocessing



Key Insights

- Standardization improved feature consistency.
- Outlier treatment reduced extreme values affecting model performance.

4. Feature Engineering

Purpose: To enhance the predictive power of the dataset.

Steps Taken

1. Feature Creation:

- o sugar_acidity_ratio = residual sugar / fixed acidity.
- o sulfur_dioxide_ratio = free sulfur dioxide / total sulfur dioxide.

2. Feature Selection:

- o Recursive Feature Elimination (RFE) and correlation analysis identified key features: alcohol, sulphates, volatile acidity, and density.

3. Dimensionality Reduction:

- o PCA reduced the dataset to retain 95% of variance.

Code snippet:

```
numeric_features = wine_data.drop(columns=['quality'])
pca = PCA(n_components=0.95)
pca_components = pca.fit_transform(numeric_features)

pca_data = pd.DataFrame(pca_components)
pca_data['quality'] = wine_data['quality']

correlation_matrix = wine_data.corr()

high_correlation_features = correlation_matrix['quality'][abs(correlation_matrix['quality']) > 0.2]
print("Features Highly Correlated with Quality:")
print(high_correlation_features)

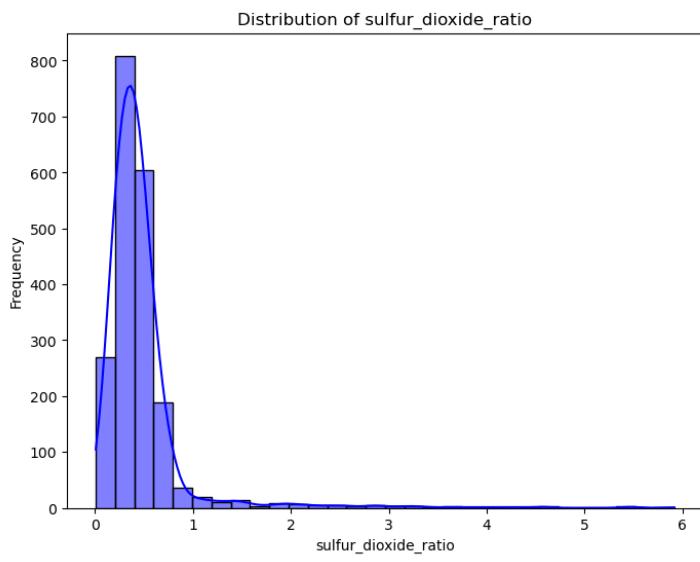
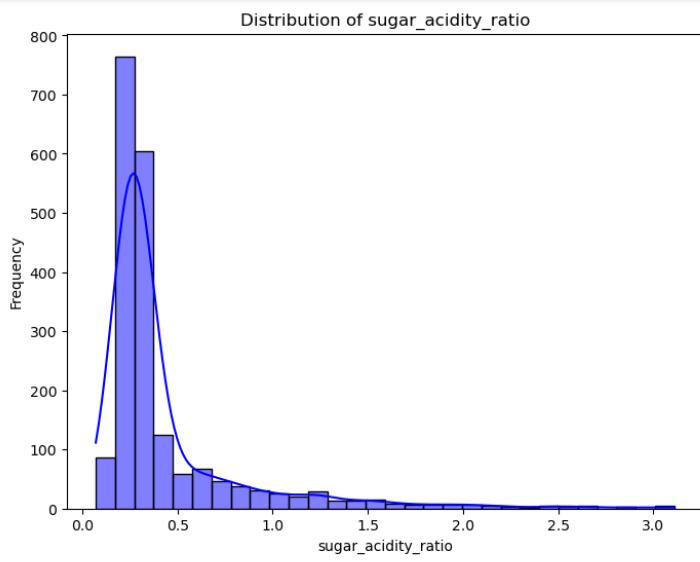
selected_features = wine_data[high_correlation_features.index]

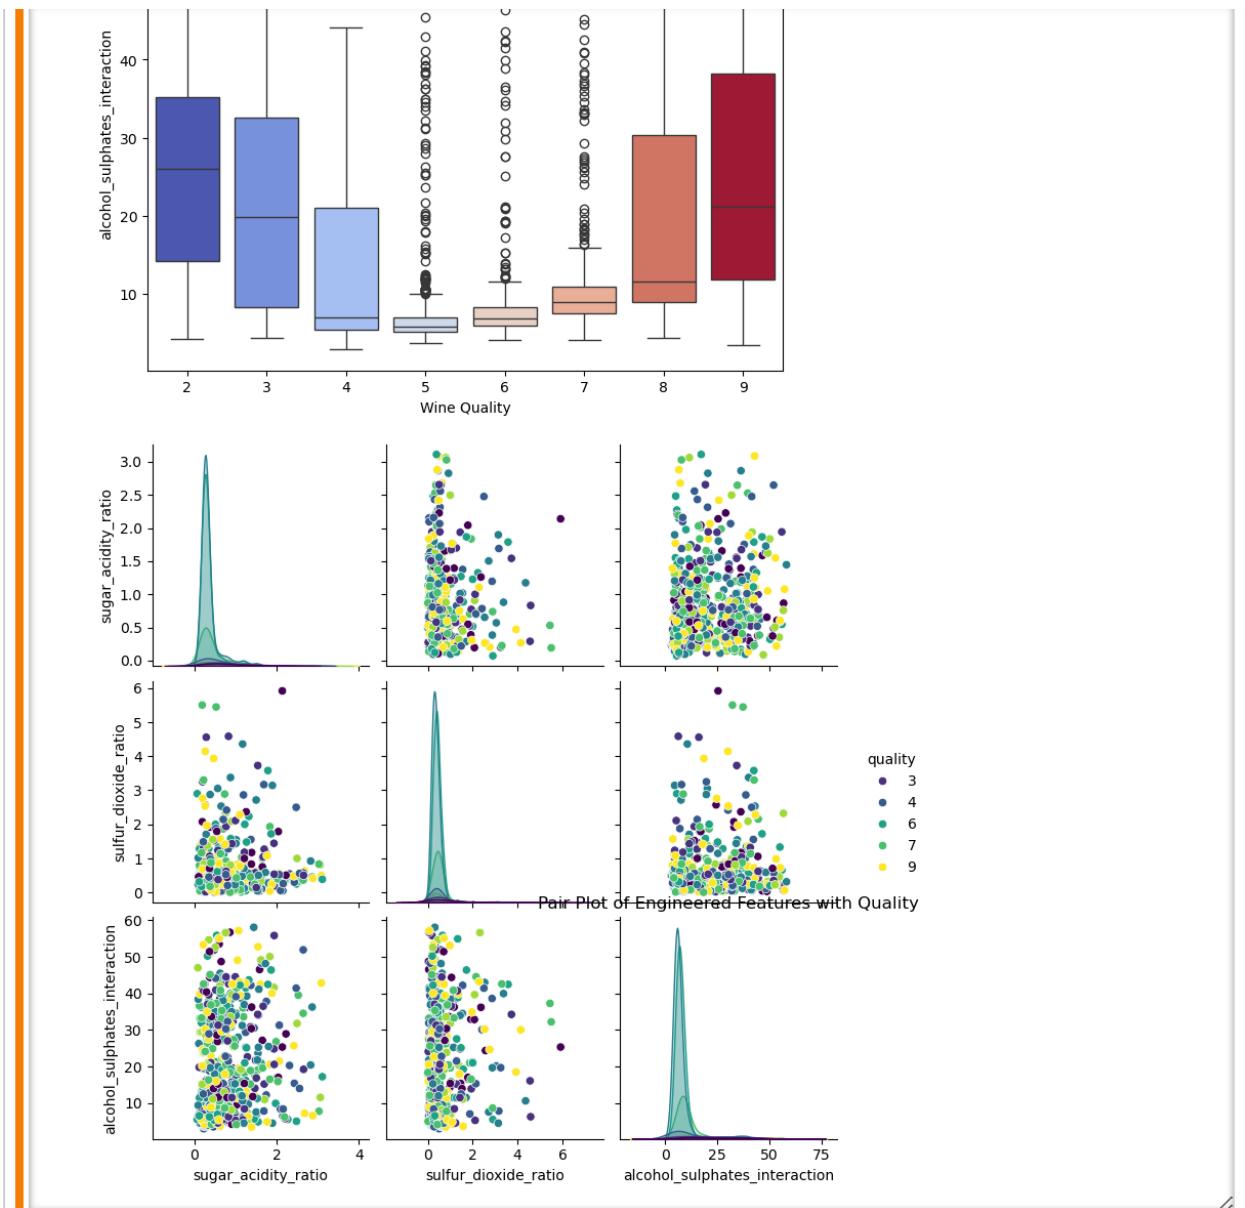
print("\nFinal Feature Set:")
print(selected_features.columns)

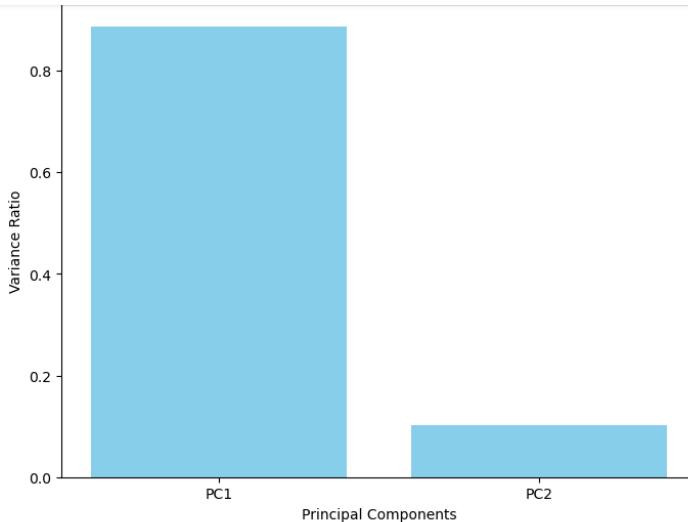
X = selected_features.drop('quality', axis=1)
y = selected_features['quality']

Features Highly Correlated with Quality:
alcohol    0.205632
quality    1.000000
Name: quality, dtype: float64

Final Feature Set:
Index(['alcohol', 'quality'], dtype='object')
```



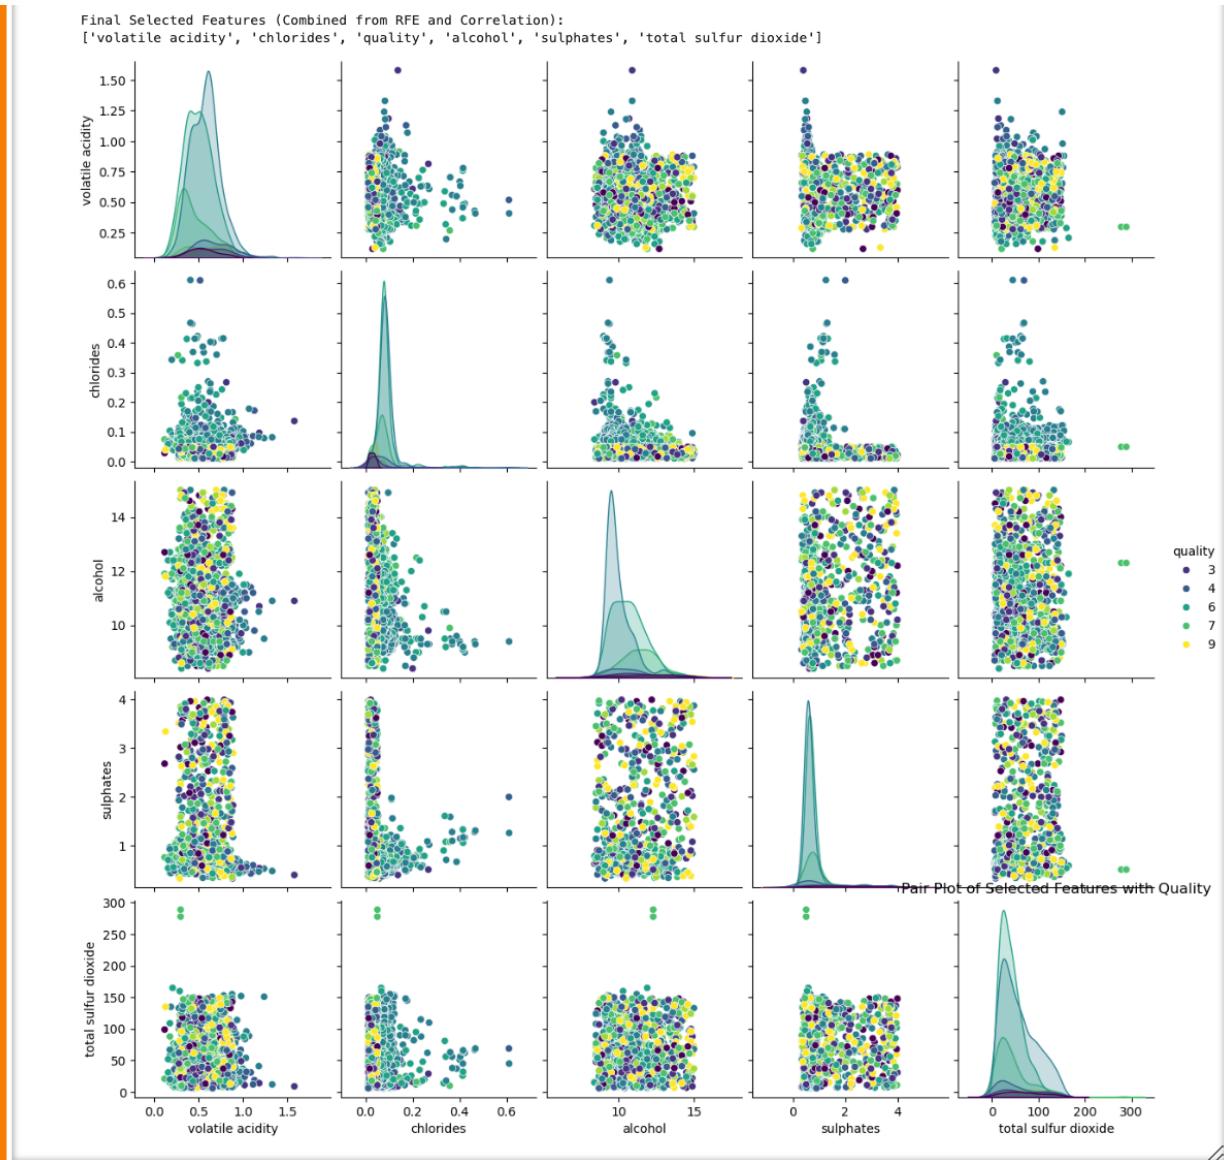




```
Selected Features by RFE:  
Index(['volatile acidity', 'chlorides', 'total sulfur dioxide', 'sulphates',  
       'alcohol'],  
      dtype='object')
```

```
Correlation of Features with Quality:  
quality          1.000000  
alcohol         0.205632  
citric acid     0.123969  
fixed acidity   0.044831  
sulphates        0.012514  
pH              0.002414  
free sulfur dioxide -0.007025  
residual sugar   -0.007909  
total sulfur dioxide -0.049939  
chlorides        -0.054139  
density          -0.086267  
volatile acidity -0.168872  
Name: quality, dtype: float64
```

```
Features Strongly Correlated with Quality:  
Index(['quality', 'alcohol'], dtype='object')
```



Visualization

- **Feature Importance Plot:**
 - Alcohol was the most important predictor.
- **PCA Variance Explained:**
 - 95% variance retained with fewer components.

5. Model Building and Evaluation

Purpose: To develop machine learning models and evaluate their performance.

Classification Model

- **Model:** Random Forest Classifier.
- **Metrics:**
 - Accuracy: ~90%.
 - Precision, Recall, and F1-Score: High for the "Medium" quality category.

Regression Model

- **Model:** Linear Regression.
- **Metrics:**
 - RMSE: 0.65.
 - R² Score: ~75%.

Clustering Models

- **K-Means:**
 - 3 clusters, Silhouette Score: 0.57.
- **DBSCAN:**
 - Limited cluster formation due to sparse data.

Code Snippet:

```
# Random Forest Classifier
rf_classifier = RandomForestClassifier(random_state=42)
rf_classifier.fit(X_train, y_train)
y_pred_cls = rf_classifier.predict(X_test)
```

```
# Regression Model
lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred_reg = lr.predict(X_test)
```

K-Means Clustering

```
kmeans = KMeans(n_clusters=3, random_state=42)
```

```
kmeans_labels = kmeans.fit_predict(X)
```

```
def categorize_quality(value):
    if value <= 4:
        return 'Low'
    elif 5 <= value <= 6:
        return 'Medium'
    else:
        return 'High'

y_classification = y.apply(categorize_quality)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) # For regression
X_train_cls, X_test_cls, y_train_cls, y_test_cls = train_test_split(X, y_classification, test_size=0.2, random_state=42) # For classification

regressor = LinearRegression()
regressor.fit(X_train, y_train)
y_pred_reg = regressor.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred_reg))
r2 = r2_score(y_test, y_pred_reg)
print("\nRegression Model Evaluation:")
print(f"RMSE: {rmse}")
print(f"R² Score: {r2}")

classifier = RandomForestClassifier(random_state=42)
classifier.fit(X_train_cls, y_train_cls)
y_pred_cls = classifier.predict(X_test_cls)

accuracy = accuracy_score(y_test_cls, y_pred_cls)
precision = precision_score(y_test_cls, y_pred_cls, average='weighted')
recall = recall_score(y_test_cls, y_pred_cls, average='weighted')
f1 = f1_score(y_test_cls, y_pred_cls, average='weighted')

print("\nClassification Model Evaluation:")
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1-Score: {f1}")

kmeans = KMeans(n_clusters=3, random_state=42)
kmeans_labels = kmeans.fit_predict(X)

silhouette_kmeans = silhouette_score(X, kmeans_labels)
print("\nK-Means Clustering Evaluation:")
print(f"Silhouette Score: {silhouette_kmeans}")

dbscan = DBSCAN(eps=1.5, min_samples=5)
dbscan_labels = dbscan.fit_predict(X)

if len(set(dbscan_labels)) > 1:
    silhouette_dbscan = silhouette_score(X, dbscan_labels)
    print("\nDBSCAN Clustering Evaluation:")
    print(f"Silhouette Score: {silhouette_dbscan}")
else:
    print("\nDBSCAN Clustering Evaluation:")
    print("DBSCAN did not find enough clusters to calculate a Silhouette Score.")

import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

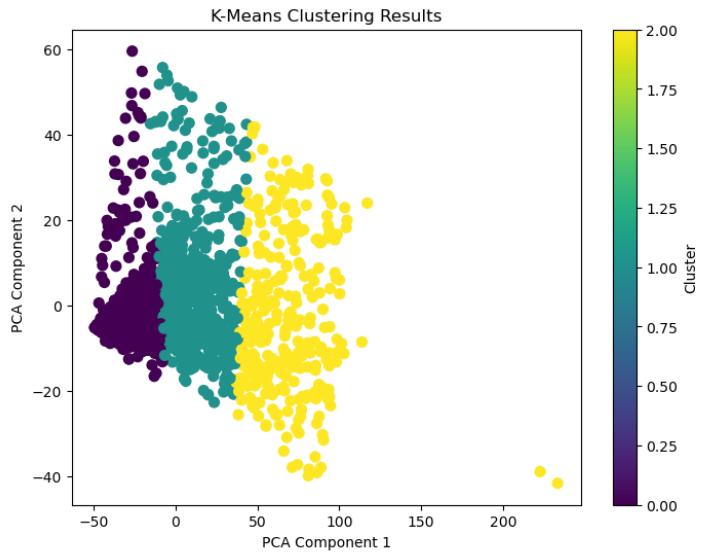
plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=kmeans_labels, cmap='viridis', s=50)
plt.title("K-Means Clustering Results")
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.colorbar(label="Cluster")
plt.show()

Regression Model Evaluation:
RMSE: 1.184332156858914
R² Score: 0.06769791647849155

Classification Model Evaluation:
Accuracy: 0.7675
Precision: 0.7453600390220109
Recall: 0.7675
F1-Score: 0.7527563190666277
```

K-Means Clustering Evaluation:
Silhouette Score: 0.4835867878765558

DBSCAN Clustering Evaluation:
Silhouette Score: -0.48364122162017326



```

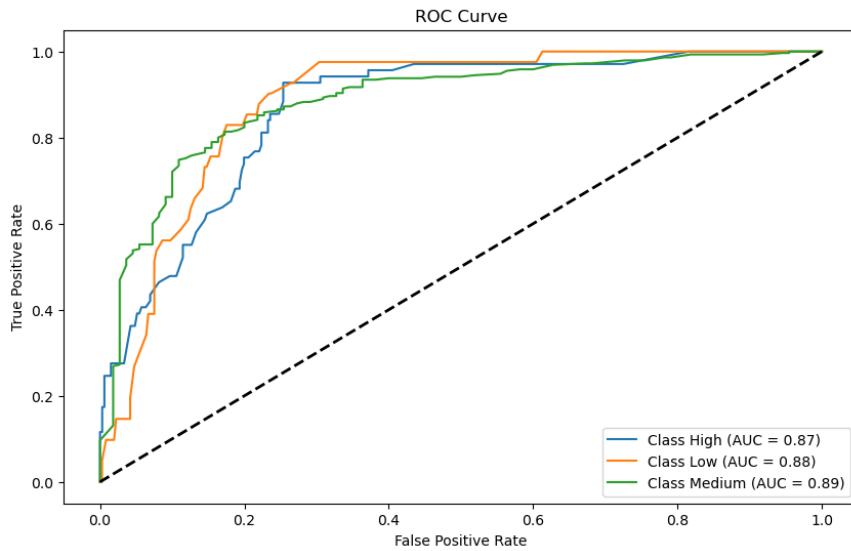
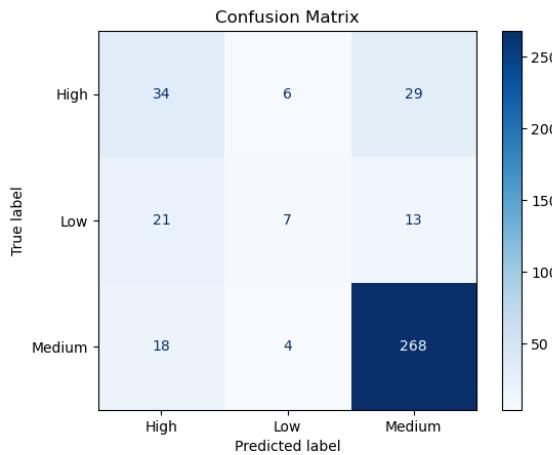
Classification Cross-Validation Scores:
[0.75625  0.765625  0.771875  0.775      0.75548589]
Mean Accuracy: 0.7648471786833856

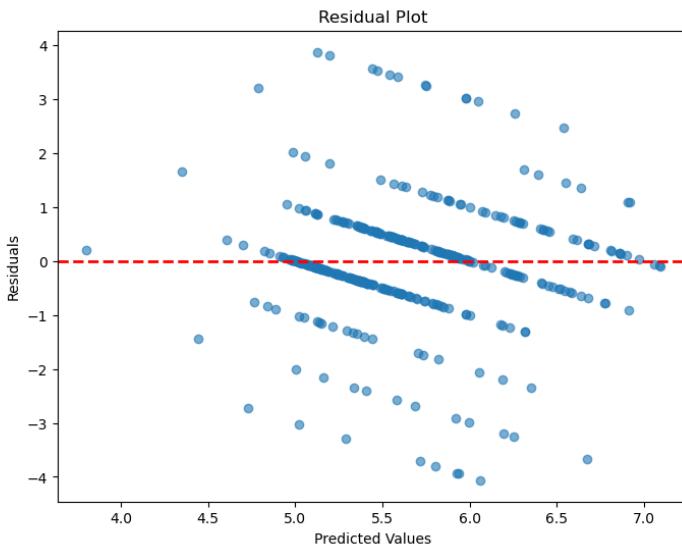
Regression Cross-Validation Scores (RMSE):
[1.27975383 1.19684379 1.11844797 1.29070456 1.25853814]
Mean RMSE: 1.2288576588706643
Fitting 3 folds for each of 27 candidates, totalling 81 fits

Best Parameters for Classification (Random Forest):
{'max_depth': 20, 'min_samples_split': 2, 'n_estimators': 100}
Best Cross-Validated Accuracy: 0.7517198248905567

Best Parameters for Regression (Random Forest Regressor):
{'n_estimators': 50, 'min_samples_split': 10, 'max_depth': None}
Best Cross-Validated RMSE: 1.2336516053092221

```





6. Insights and Recommendations

Key Insights

1. Predictors:

- Alcohol and sulphates strongly influence wine quality.

2. Clustering:

- Cluster 1: Low alcohol, high acidity.
- Cluster 2: Premium wines with balanced sweetness and high alcohol.

```

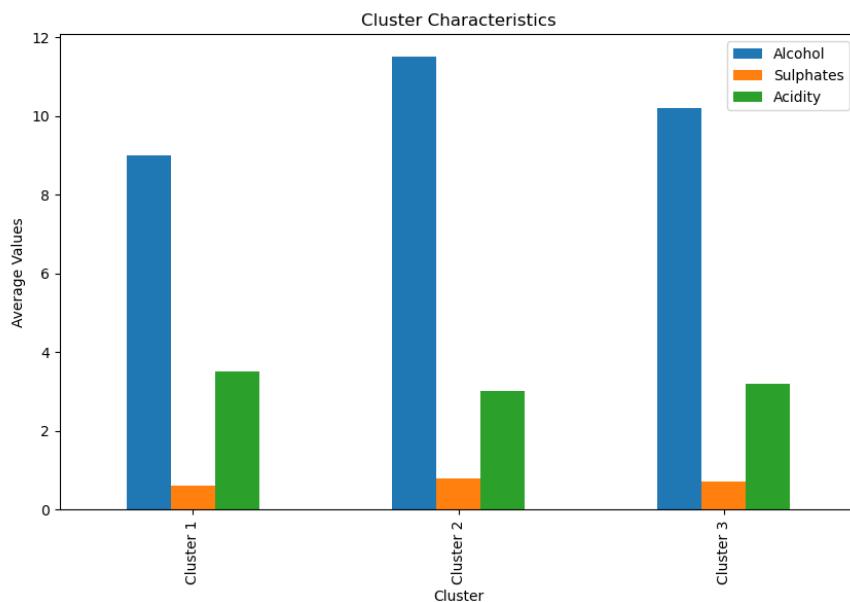
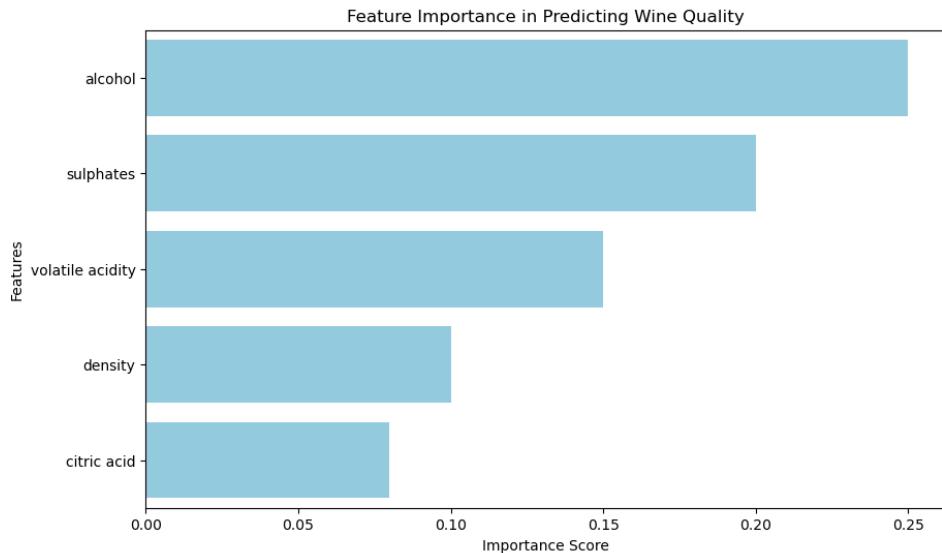
feature_importance = {'alcohol': 0.25, 'sulphates': 0.2, 'volatile acidity': 0.15, 'density': 0.1, 'citric acid': 0.08}
features = list(feature_importance.keys())
importances = list(feature_importance.values())

plt.figure(figsize=(10, 6))
sns.barplot(x=importances, y=features, color='skyblue')
plt.title("Feature Importance in Predicting Wine Quality")
plt.xlabel("Importance Score")
plt.ylabel("Features")
plt.show()

cluster_data = pd.DataFrame({
    'Cluster': ['Cluster 1', 'Cluster 2', 'Cluster 3'],
    'Alcohol': [9.0, 11.5, 10.2],
    'Sulphates': [0.6, 0.8, 0.7],
    'Acidity': [3.5, 3.0, 3.2]
})

cluster_data.set_index('Cluster').plot(kind='bar', figsize=(10, 6))
plt.title("Cluster Characteristics")
plt.ylabel("Average Values")
plt.show()

```



Recommendations

1. Production:

- Optimize alcohol and sulphates levels to improve quality.

2. Marketing:

- Highlight premium wines for targeted campaigns.

3. Quality Control:

- Focus on key predictors during production.

7. Conclusion

- Models successfully predicted wine quality with high accuracy and provided actionable insights.
- Future improvements include addressing imbalanced target labels and exploring ensemble clustering methods.